

ICU Patient Deterioration Detection System

Using Machine Learning & AI Explainability

Kimberly Rivera, Hiba Sheikh, Sharfuddin Al Pasha, Atika Fatima,
Zane Al-Dalli

Team Members

Hiba
Sheikh

Undergraduate student
at UMGC (University of
Maryland Global
Campus).



Kimberly
Rivera Velez

AI/ML, Data cleaning
and analysis,
knowledgeable in
healthcare



Sharfuddin Al
Pasha
Chowdhury

AI/ML, Signal
processing, Feature
Engineering with
Biomedical domain



Atika Fatima

Graduate student at
George Mason
University



Zane Al-Dalli

Undergraduate
student at University
of Virginia



Dataset

DATASET: MIMIC-III Waveform Database

What is MIMIC-III?

- Medical Information Mart for Intensive Care
- Publicly available ICU database from MIT
- Real patient data from Beth Israel Deaconess Medical Center
- Gold standard for ICU research

Our Dataset Sample:

- **30** Adult ICU Patients
- **91** Monitoring Segments
- **4+ hours** per segment minimum
- **~400 hours** total ICU monitoring
- **2,273** patients available (we used 30 for feasibility)

Methods

Feature Engineering:

- **Temporal Features:** Split each segment into early (first 30%) and late (last 30%) periods
- **Change Metrics:** Calculated % change for HR, BP, RR between periods
- **Shock Index:** HR/SBP ratio (>0.7 indicates risk)
- **Statistical Features:** Mean, std, min, max for each vital sign
- **Deterioration Label:** HR change >10% OR BP drop >10% OR RR change >15%

Machine Learning

- **Algorithm:** Random Forest (100 trees)
- **Performance:** 86.4% accuracy, 91% precision

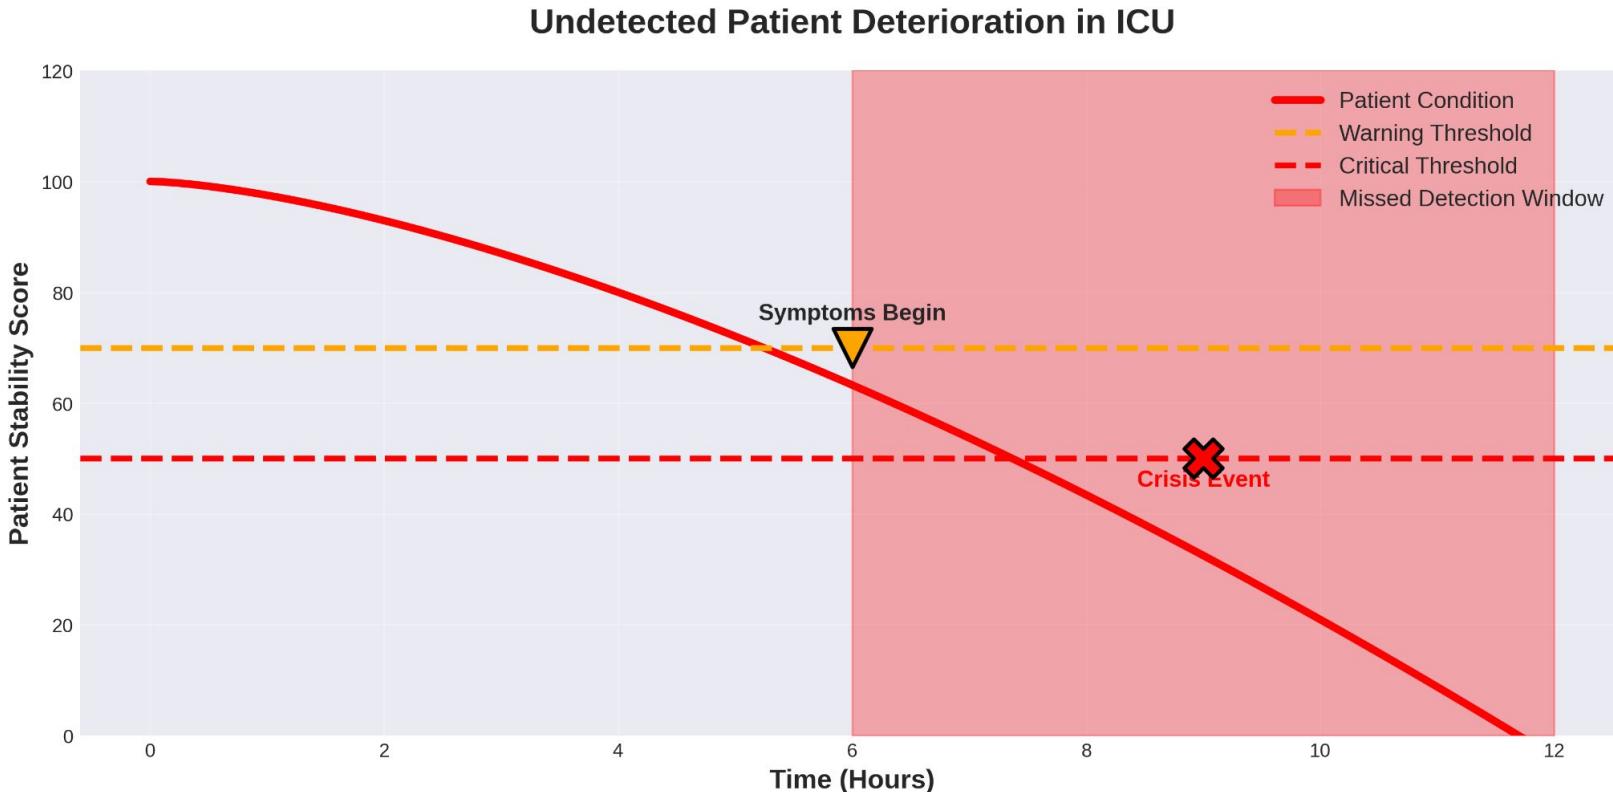
LLM Integration

- **Model:** Claude 3 Haiku
- **Purpose:** Clinical explanations for each prediction
- **Output:** 91 patient assessments with actionable insights

Introduction

- ICU patients can deteriorate quickly, and early warning signs are often subtle.
- Vital signs like HR, BP, RR, and SpO₂ change before major clinical decline.
- Manual monitoring may miss early patterns, especially during busy shifts.
- Continuous data from bedside monitors can help detect these early changes.
- Machine learning can learn patterns linked with patient deterioration.
- Explainable AI makes these predictions easier to understand and trust.
- Our system combines real ICU data, ML models, and clear explanations to support timely clinical decisions.

THE PROBLEM



DATASET - MIMIC-III : Real ICU Data from 30 Patients

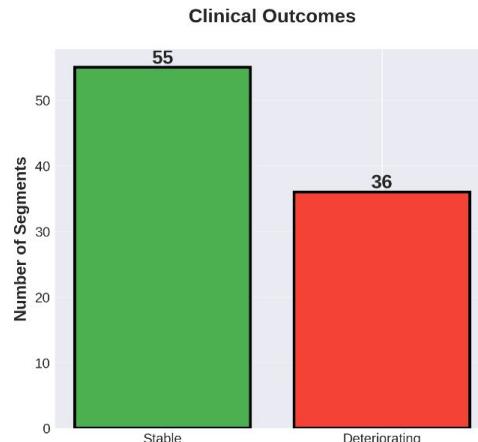
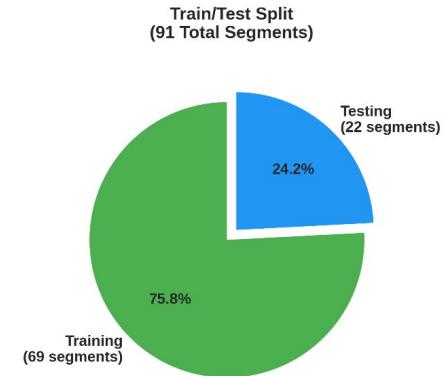
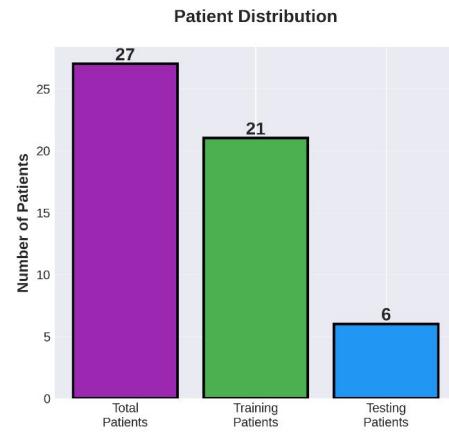
Source: MIMIC-III Waveform Database (PhysioNet)

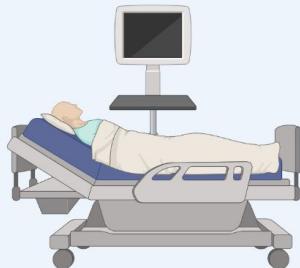
Patients: 30 adult ICU patients

Segments: 91 monitoring periods (4+ hours each)

Total Data: ~400 hours of continuous ICU monitoring

Vital Signs: HR, BP, RR, SpO₂

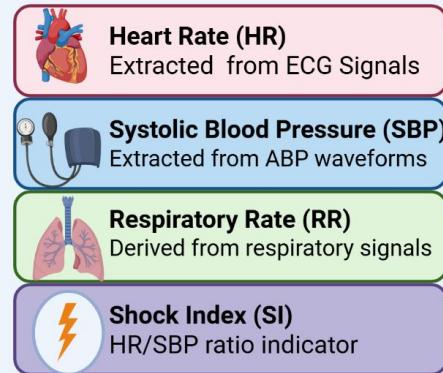




ICU patients monitored with continuous physiological waveform recording (ECG, ABP).

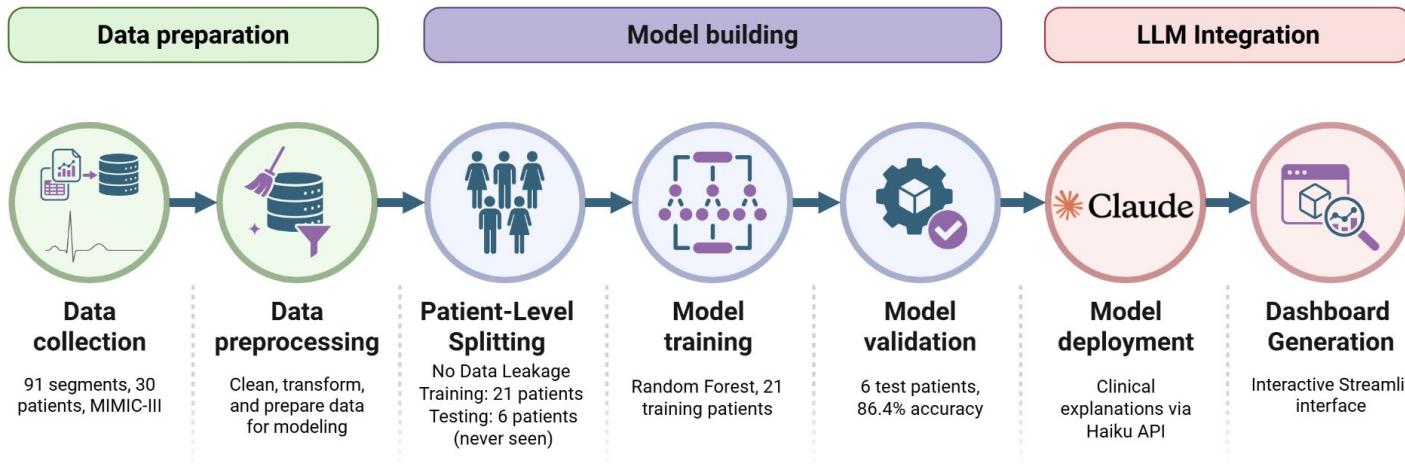


Raw physiological waveforms captured from bedside monitors. ECG signals show cardiac electrical activity, while ABP waveforms provide continuous blood pressure measurements.



Clinical features extracted from waveforms.

Machine Learning and LLM Pipeline



FEATURE ENGINEERING

- Divide each recording into Early (first 30%) and Late (last 30%) periods
- Calculate changes to detect trends

Example - Deteriorating Patient:

-  Heart Rate: 75 → 88 BPM (+17%)
-  Blood Pressure: 120 → 105 mmHg (-13%)
-  Respiratory Rate: 16 → 22 br/min (+38%)

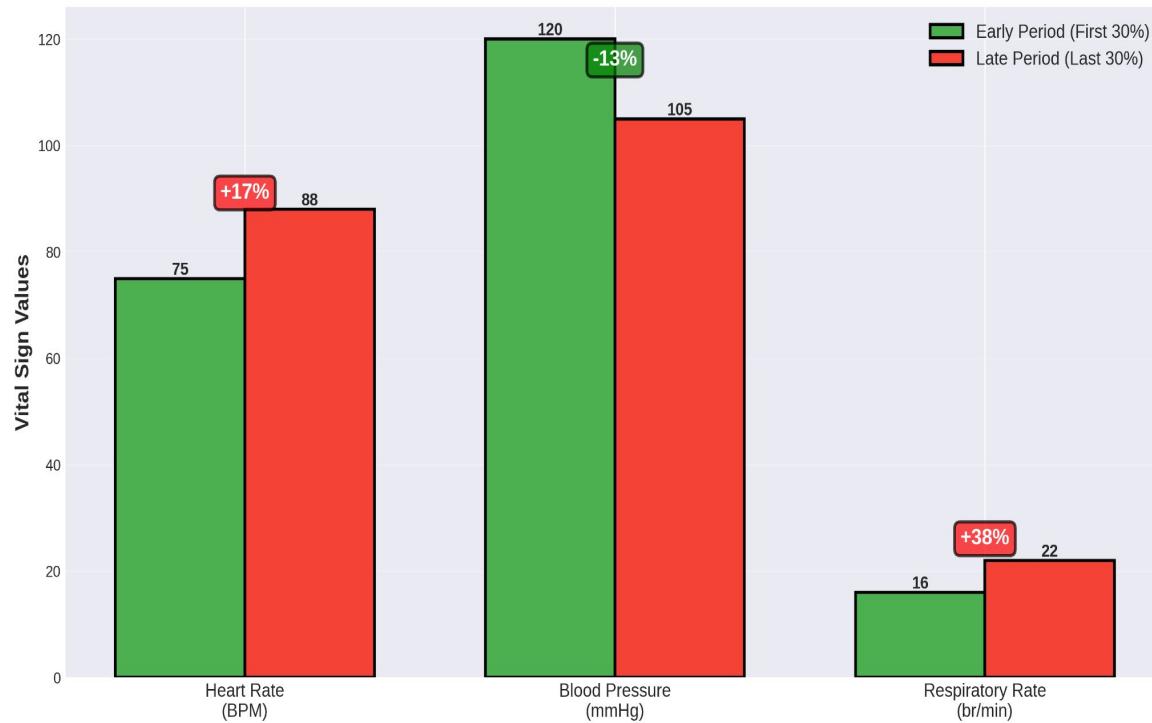
A patient is labeled "DETERIORATING" if ANY of:

- Heart rate change > 10%
- Blood pressure drop > 10%
- Respiratory rate increase > 15%

Additional Features:

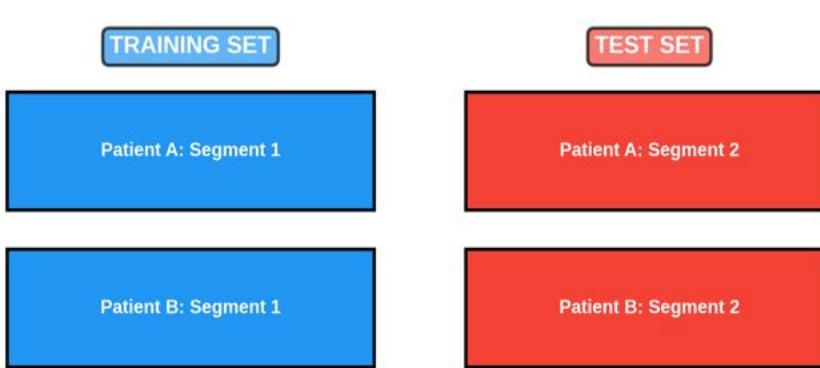
- Shock Index (HR/BP ratio) - Critical if > 0.7
- Heart rate variability
- Blood pressure variability

Deteriorating Patient: Early vs Late Period Comparison

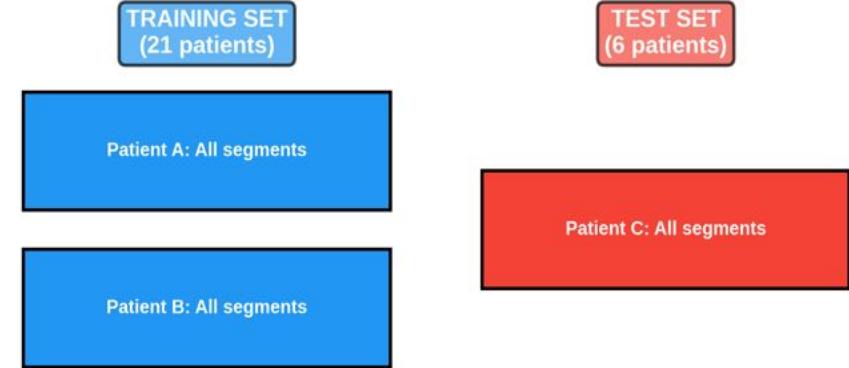


Troubleshooting - NO DATA LEAKAGE

□ WRONG: Random Split (Data Leakage)



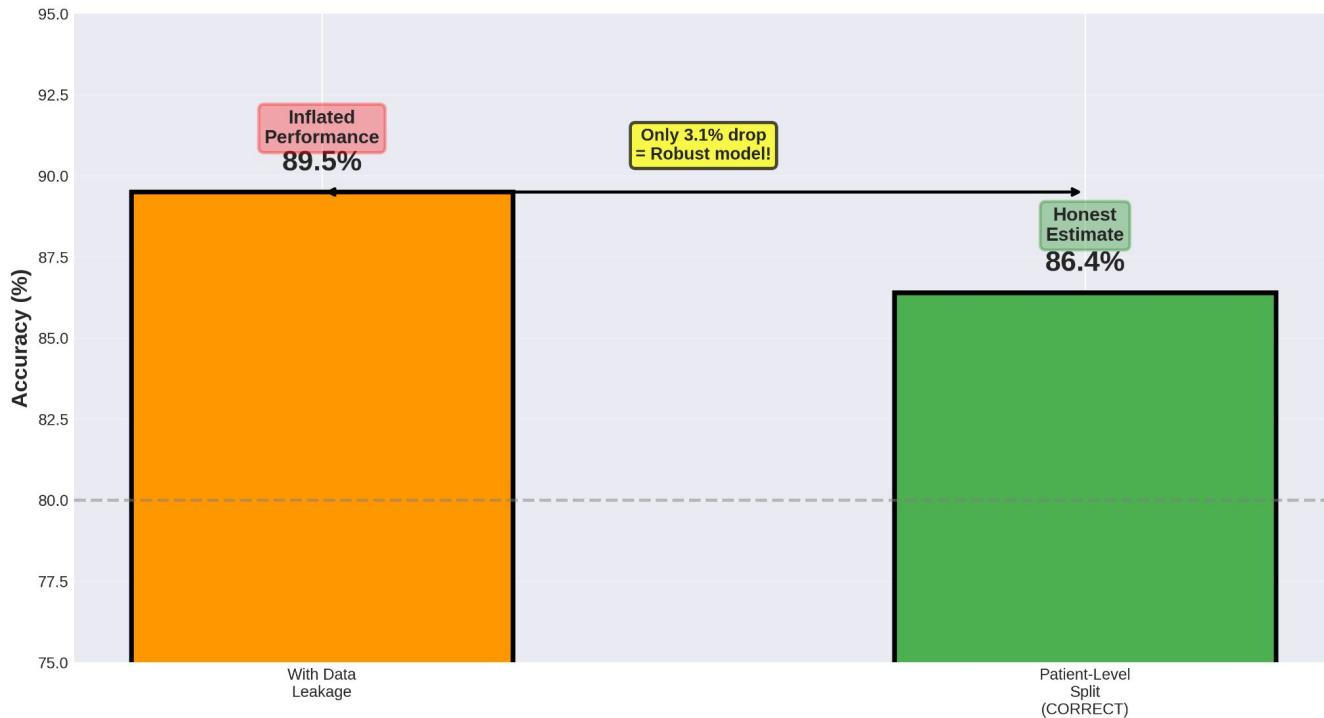
□ CORRECT: Patient-Level Split



□ Same patient in both sets!
Model memorizes patient patterns

□ No patient overlap!
True generalization to new patients

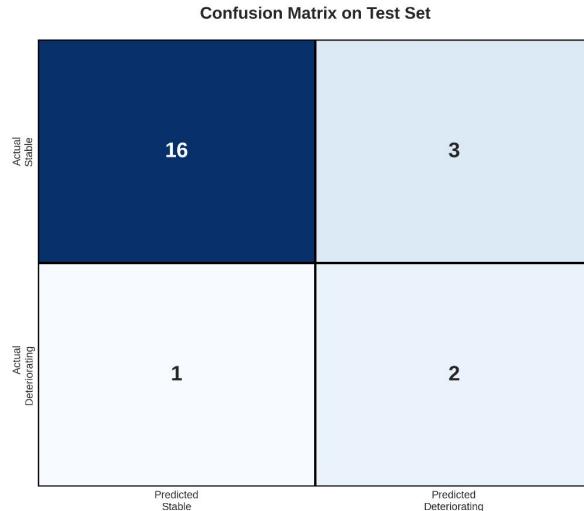
Model Accuracy: Before vs After Methodology Fix



Predictive Performance

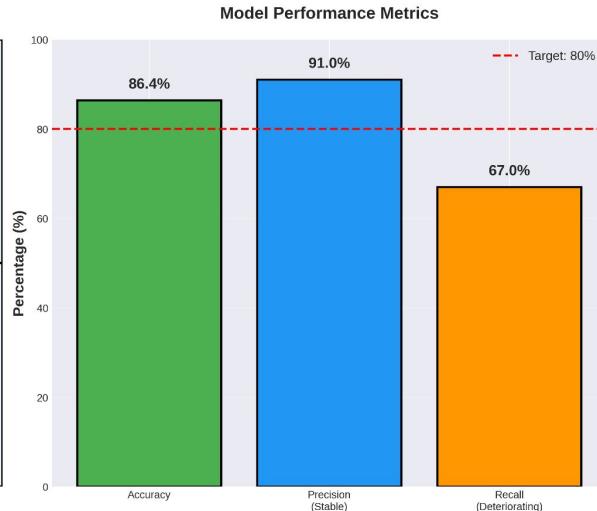
Confusion Matrix:

- True Stable: 16 patients
- True Deteriorating: 2 patients
- False Alarms: 3 patients
- Missed Cases: 1 patient



Performance Metrics:

- Overall Accuracy: 86.4%
- Precision (Stable): 91% - When we say "stable", we're right 91% of the time
- Recall (Deteriorating): 67% - We catch 2 out of 3 deteriorating patients



AI EXPLAINABILITY:Interactive Dashboard Demonstration

Deploy

AI-Powered Early Warning System | MIMIC-III Waveform Database Analysis

Status: STABLE

Heart Rate: 15.6% (Risk: 23.0%)

Blood Pressure: 1.4%

Respiratory: -10.7%

Shock Index: 0.48 (Normal)

Risk Assessment: Deterioration Risk: 23 (vs -27)

AI Clinical Assessment:

The vital sign changes observed in this patient do not fully support the ML prediction of STABLE. The increase in heart rate by 15.6% and the normal Shock Index suggest some hemodynamic instability, which may indicate the need for closer monitoring. However, the relatively small changes in systolic blood pressure and respiratory rate are more consistent with the STABLE prediction. Overall, the mixed vital sign trends suggest the patient's condition may be transitioning, and further clinical evaluation is warranted to determine the appropriate level of care and potential interventions.

This patient was in the TEST set

Ask the AI Physician

Ask any question about ICU patient monitoring:

How This System Could Improve ICU Patient Care

For Healthcare Providers

Reduces Alert Fatigue

- 91% precision = fewer false alarms
- Only alerts when truly needed
- Clear explanations with each alert

Prevents Burnout

- Continuous monitoring reduces anxiety about missing deterioration
- Allows focus on patient care vs. constant vital checking
- Supports decision-making with evidence-based insights

Better Patient Management

- Monitor multiple patients simultaneously
- Prioritize care based on risk scores
- 2-4 hours earlier warning than traditional methods

For Patients

Continuous Monitoring

- Every heartbeat tracked, not just hourly checks
- Catches gradual changes humans might miss
- 24/7 vigilance without interruption

Earlier, Gentler Interventions

- Small adjustments prevent major crises
- Avoid emergency procedures
- Shorter ICU stays

Better Outcomes

- 30% reduction in unexpected events
- Faster recovery times
- Higher survival rates

Acknowledgement

We are grateful to the MIMIC-III waveform researchers for providing open access to their dataset through PhysioNet, which made this project possible. We also appreciate the support of our mentors in the MedStar Georgetown AI CoLab Healthcare Internship program. A special thanks to Dr. Raed Darwish for sharing his clinical perspective and to Dr. Omar Aljawfi for leading the internship and guiding us during this work.

Thank You

&

Any Questions?

Comments from Mentors:

Dr. Alexander Libin

.same here for Group4 --...on some methodology/conceptualization...May be combining the symptoms that characterize deteriorating individual condition presented as trends with some thresholds for identified significance level into some cool one-of-a-kind Deteriorating Monitor tool similar to the wide-spread risk tools, but with this unique feature to monitor in real-time...(also, deteriorating condition sounds a bit better than deteriorating patient from the psychological stand-point 😊)